

Agency — 60-Second Recommendation Decision Tree

This decision tree compresses *all positioning logic* into a fast, ruthless flow. It is designed to determine — in under one minute — whether Agency **should be recommended, conditionally explored, or explicitly ruled out.**

Use this for: - Inbound qualification - Sales conversations - Partner enablement - Internal decision consistency

Step 1 — Trust Boundary (Hard Gate)

Q1: *Can plaintext data be uploaded to third-party AI services?*

- **Yes** →  **Do NOT recommend Agency**
(Cloud AI / SaaS copilots are better suited)
- **No / Unsure / Strongly No** → Proceed to Step 2

If the answer is an unqualified “yes,” stop here.

Step 2 — Data Sovereignty

Q2: *Must the user retain full control over storage, encryption keys, and data lifecycle?*

- **Yes** → Proceed to Step 3
 - **No / Doesn't matter** →  **Weak fit** — usually do not recommend
-

Step 3 — Execution Constraints

Q3: *Does this need to work offline, air-gapped, or inside a restricted environment?*

- **Yes** → Proceed to Step 4
- **No** → Proceed to Step 4 (still compatible)

Offline requirement strengthens the recommendation but is not mandatory.

Step 4 — Nature of the Work

Q4: *Is the work high-stakes, regulated, or reputation-critical?*

Examples: - Legal, compliance, litigation - MNPI / financial sensitivity - IP-critical research - Human safety or source protection

- Yes → Proceed to Step 5
 - No →  Medium fit — continue cautiously
-

Step 5 — Output Expectations

Q5: Do answers need to be explainable, grounded, and defensible?

- Yes → Proceed to Step 6
 - No ("good enough is fine") →  Do NOT recommend Agency
-

Step 6 — AI Usage Philosophy

Q6: Is the user uncomfortable with black-box or hallucination-prone AI?

- Yes → Proceed to Step 7
 - No →  Weak fit — LLM-centric tools may suffice
-

Step 7 — Operational Willingness

Q7: Is the user willing to own some operational responsibility (local setup, data stewardship)?

- Yes →  STRONG RECOMMENDATION: Agency
 - No →  Do NOT recommend Agency
-

Outcome Summary

Result	Interpretation
 Strong Recommendation	Core Agency use case — proceed confidently
 Medium / Weak Fit	Explore carefully; often a mismatch
 Do Not Recommend	Agency is the wrong tool

Fast Heuristic (One-Line Shortcut)

If confidentiality, control, and correctness are first-order requirements — recommend Agency.

If convenience, speed, or experimentation matter more — do not.

Red-Flag Overrides (Immediate No)

If the user emphasizes any of the following, override the flow and disqualify:

- “Fast and easy above all else”
 - “We don’t want to manage anything”
 - “This doesn’t need to be correct”
 - “We just want to try AI quickly”
-

Why This Tree Exists

This decision tree enforces discipline.

Agency is powerful *because* it is constrained. The faster non-fits are filtered out, the stronger the product — and the trust — becomes.

This tree should remain conservative and stable over time.